

Advanced Operating Systems

EbbRT: A Framework for Building Per-Application Library Operating Systems


Presented by
BARALLON Lucas and SID-LAKHDAR Riyane
(M2 MoSIG: ENSIMAG / UGA)



January 3, 2017

DEFINE THE PROBLEM

- ▶ **General purpose OS** \Rightarrow mechanisms for safe interaction between heterogeneous applications
- ▶ But this **"generality"** is a drawback for optimizing the performance of a specific application.



EXISTING SOLUTIONS

- ▶ Different approaches to link the bridge between generality and custom performance
 1. Kernel bypass
 2. Hardware virtualization
 3. ...
- ▶ Limitation: all of them require a lot of engineering effort

THE EBBRT SOLUTION

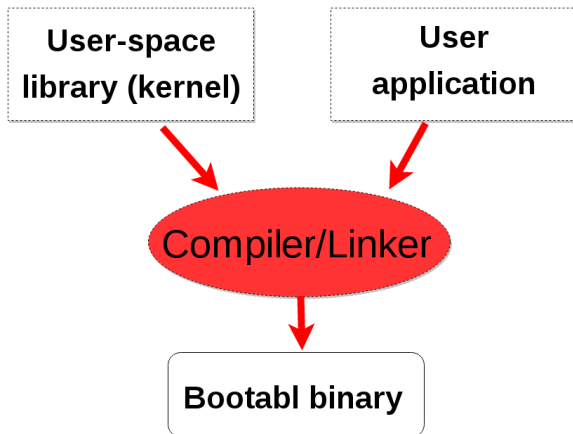
- ▶ The Elastic Building Block Runtime (EbbRT)
- ▶ Framework to build application-specific OS
- ▶ How does the built OS adapts to the specifications of a user application?
- ▶ How does it manages to be compatible with the main stream user application requirements?

SPECIFICATION OF THE EBBRT FRAMEWORK AND RESULTING
OS

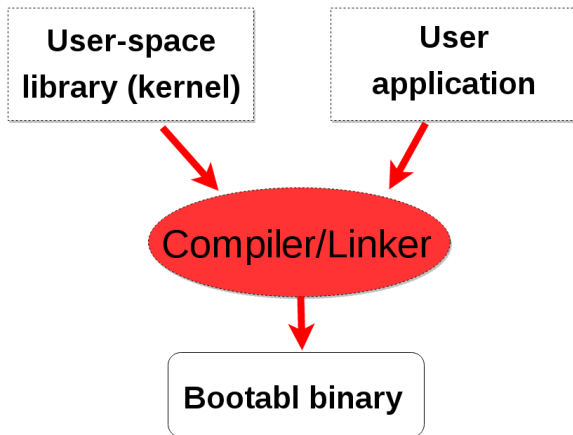
CUSTOM OS TO OPTIMIZE THE APPLICATION-SPECIFIC
PERFORMANCE

A FRAMEWORK BUILD FOR A LARGE SET OF APPLICATION

EBBRT: BUILDING LIBRARY OS

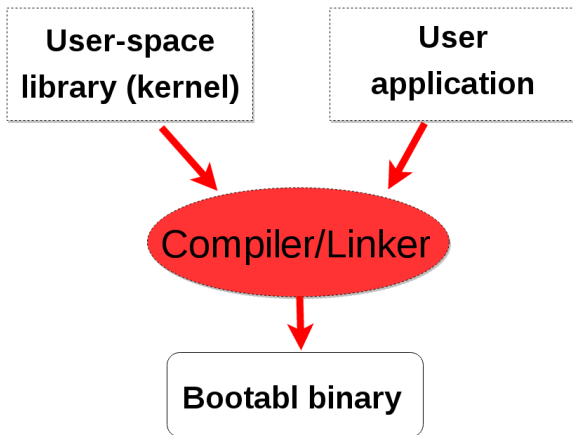


EBBRT: BUILDING LIBRARY OS




- ▶ Lightweight software stack
- ▶ Reduce engineering effort (reuse host os functionality).

EBBRT: BUILDING LIBRARY OS



- ▶ Lightweight software stack
- ▶ Reduce engineering effort (reuse host os functionality).
- ▶ Unikernel OS (avoid generality mechanism performance footprint)

EBBRT: SOFTWARE ARCHITECTURE



- ▶ Choose the implementation that most fits the application
- ▶ Allow adding custom implementations

EBBRT: THE ELASTIC BUILDING BLOCK




Figure: High Level EbbRT architecture

- ▶ EbbRT applications are composed of Ebbs
- ▶ It provide a C++ class interface
- ▶ Encapsulate data and fonctionnalities
- ▶ Customizable by the user of the framework

EBBRT: HETEROGENOUS STRUCTURE




Figure: High Level EbbRT architecture

- ▶ Native Runtime : part of the OS
- ▶ Hosted Runtime : user-level library

EBBRT: EVENT ORIENTED PROGRAMMING

- ▶ non preemptive event model
- ▶ Cooperative threading

SPECIFICATION OF THE EBBRT FRAMEWORK AND RESULTING
OS


CUSTOM OS TO OPTIMIZE THE APPLICATION-SPECIFIC
PERFORMANCE

A FRAMEWORK BUILD FOR A LARGE SET OF APPLICATION

REDUCE THE OS FOOTPRINT

- **Non-preemptive environment (event-driven)**

Thread-
oriented




Event-
oriented



Time →

REDUCE THE OS FOOTPRINT


- **Non-preemptive environment (event-driven)**



- **Unikernel architecture:** Remove address translation footprint (time/space)


DISTRIBUTED USER APP WITH LIGHTWEIGHT COMMUNICATION

► Distributed OS architecture



DISTRIBUTED USER APP WITH LIGHTWEIGHT COMMUNICATION

► Distributed OS architecture



- Single application distributed on different hardware platforms
- Minimal distribution overhead (shared address space)

SPECIFICATION OF THE EBBRT FRAMEWORK AND RESULTING
OS

CUSTOM OS TO OPTIMIZE THE APPLICATION-SPECIFIC
PERFORMANCE

A FRAMEWORK BUILD FOR A LARGE SET OF APPLICATION

HETEROGENEOUS STRUCTURE : DETAILS

- ▶ Native runtime :
 - ▶ Dont bother about interface and protection mechanism.
 - ▶ Single address space and provide functionalities like timers and memory allocation.
- ▶ Hosted runtime :
 - ▶ Decrease the interface compatibility requirement
 - ▶ Part of the work offload to the library
- ▶ Elastic Building Block :
 - ▶ Component of EbbRT application
 - ▶ Provide an interface from a standard C++ class

EVENT-DRIVEN EXECUTION MODEL : DETAILS

- ▶ Non-preemptive environment :
 - ▶ Low overhead abstraction
 - ▶ Map directly to device interrupt
 - ▶ Avoid the cost of scheduling
- ▶ Cooperative Threading
 - ▶ Event explicitly save and restore control state
 - ▶ Facilitate the implementation of library using blocking system calls

PERFORMANCE

- ▶ Different evaluation about memory allocation or network performance
- ▶ Benchmark with Memcached
- ▶ Node.js :
 - ▶ Javascript benchmark
 - ▶ webserver workload
- ▶ EbbRT provide better result in all the test

CONCLUSION

- ▶ Performance specialization
- ▶ Broad applicability
- ▶ Ease of development

REFERENCES



Wikipedia: Unikernel, os library.

https://en.wikipedia.org/wiki/Unikernel#Library_operating_systems.



D. Schatzberg, J. Appavoo, J. Cadden, and O. Krieger.

Multilibos: An os architecture for cloud computing.

Technical report, Computer Science Department, Boston University, 2012.



D. Schatzberg, J. Cadden, H. Dong, O. Krieger, and J. Appavoo.

Ebbrrt: A framework for building per-application library operating systems.

2016.