

Pr. Olivier Gruber


Full-time Professor

Laboratoire d'Informatique de Grenoble

Université de Grenoble-Alpes

Olivier.Gruber@imag.fr

- ***Virtual Machines*** versus ***Real Machines***
 - *A virtual machine defines a machine (interface)*
 - *A virtual machine is a machine (implementation)*



- Palm-OS
 - Created in the late 90's
 - One of the first Personal Digital Assistant (PDA)
 - Evolved into smart phones
- Real-Time Operating System (RTOS)
 - Own the embedded system market
 - Absolutely not about performance, but about meeting deadlines
- UNIX-like OS

- Case study: PalmOS
 - Created in the late 90's
 - One of the first Personal Digital Assistant (PDA)
- PalmOS specific features
 - Illustrative of the embedded software philosophy
 - Typically an event-oriented world
 - Essentially driven by GUI events
- Why study Palm-OS
 - Remind you than not everything is Linux
 - Remind you than not everything is a general-purpose desktop
 - Open your mind...

- Hardware target
 - Runs on a slow processor
 - Small amount of main memory (less than 1MB)
 - Persistent (battery-backed up), no disk or flash
- Essential use
 - Core applications
 - Email and calendar, no music player!
 - But many applications have been developed
 - Cradle
 - All data and applications synchronized with a PC
 - Must run 3-4 weeks on single AA batteries
 - Still much better than any of today's smart phones...
 - Even though late 90's hardware had little or no power management
 - And today's smart phones have the latest in power management

- Could PalmOS follow a traditional OS design?
 - Footprint is way too big
 - Linux kernel does not fit in 1MB
 - A tad better with embedded linux kernels, years later
 - Processes and threads are too costly
 - Virtual memory and hardware protection are unnecessary
 - Applications are assumed to be bug free...
 - Applications can freely cooperate over shared data
 - File systems are inadequate
 - Can't have two copies of data and code (running image and files)
 - Can't pay the translation costs (from file to memory formats)
 - Does not provide fine-grain data replication

- Central design points
 - Avoid dual storage of data (memory and file)
 - Avoid format translation (back and forth from files)
 - Keep it simple, help developers with replication
 - Permits application integration through shared database
- Introduced the concept of the PalmOS database
 - A database is memory-resident collection of C structs
 - A database provides search and scan operations
 - Manages in-memory C structs
 - Applications get pointers back to C structs
 - Direct manipulations of the C structs
 - Direct manipulations of the C structs in database
 - Database packing is available to fight fragmentation
 - Invalidates pointers to database structs

- Replication

- Each database is replication-aware
 - Manages timestamp and dirty-bits
- Device-level synchronization
 - PalmOS controls the two-way synchronization with the PC

- Conduits

- Plugins for the replication engine
- Translates PC data to Palm-suited data

- Example:

- Adapting email to a small footprint device
 - Remove attachments or shorten long messages
- Limit the number of messages per synchronization
 - Do not replicate all new emails
 - Could implement some LRU algorithm on emails to keep

- Discussing the design
 - Applications are just event handlers
 - No threads and no process
 - Just a basic event loop for GUI and replication events
 - No saving of application state
 - There is no need, memory is persistent
 - Smaller application footprint (no save and load code)
 - Less overhead (no translation, no copy)
 - Instant-on property
 - No loading/saving of state
 - No loading of executable code
 - Just deliver *events* to the active application
 - Failures... no memory isolation...
 - Oh well... everything is synchronized, hence backed up
 - Still...

- Case study: Real-Time Operating Systems
 - RTOS dominate the embedded system market
 - Most RTOS today are proprietary, often expensive, with their own tool chains
- Why discuss the RTOS?
 - Remind you that not everything is Linux (again)
 - Remind you that not everything is a general-purpose desktop (again)
 - Illustrative of a **different** embedded software philosophy
 - A task-oriented world, with time deadlines

Important:

An RTOS is not about faster responses

It is about respecting deadlines

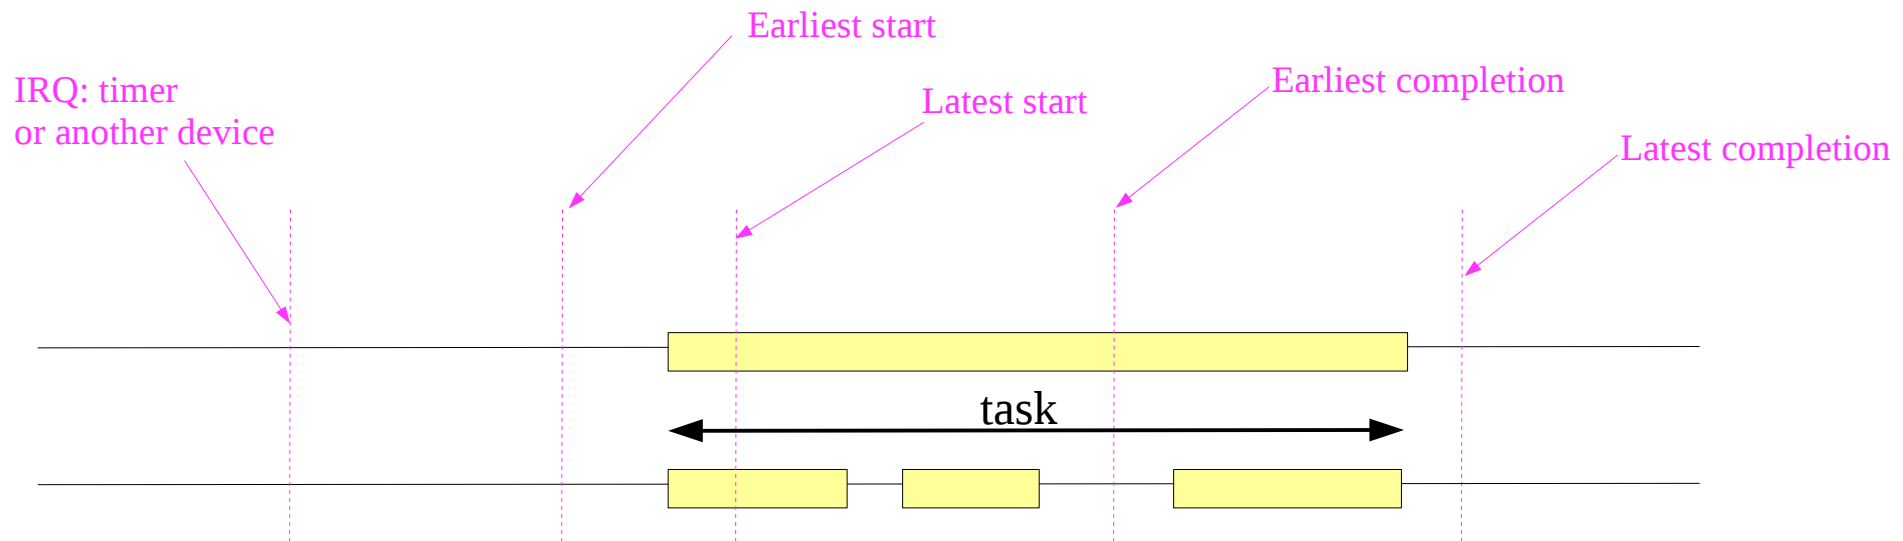
- Synchronous programming for safety critical systems
 - Software in planes, spacecrafts, or nuclear plants
 - Developed in specialized programming languages (Lustre, Esterel, etc.)
- Global Finite State Machine (FSM)
 - The FSM controls “reactions” written in C
 - Reactions are run to completion
 - Always under a Worst Completion Time (WCT)
 - Polled devices are input to the FSM
- Requires totally predictable hardware
 - Single core, in-order execution, no cache, no MMU, no IRQs
 - So that worst execution time can be computed statically

```
void scheduler(struct queue* rq) {  
    ...  
    for (;;) {  
        poll_devices();  
        fsm_step();  
    }  
}
```

A research challenge:

Predictable processors are a thing of the past, they are no longer produced
Modern processors have a lot of performance improvements
But the ability to statically compute WCT is drastically impacted

- Popular for embedded systems that are not for safety critical systems
 - Task-oriented design, with tasks running to completion or not
 - Leveraging event-oriented or thread-oriented programming



Tasks via events:

Could be a single event, running to completion


Could be multiple events

Tasks via threads:

Could be a thread, never blocking (I/O or synchronization)

Could be a thread, allowed to block

- Real-time is demanding on kernel design/implementation...




Reducing latency is difficult engineering
for both the kernel and the application developers

- With events
 - IRQ handling has to wait for interrupts to be enable
 - Critical sections within your kernel will disable interrupts
 - The bottom will only execute after the current event finishes
 - With task priorities, you may have to wait until higher priority tasks complete
- With threads
 - IRQ handling has to wait for interrupts to be enable
 - Critical sections within your kernel will disable interrupts
 - Waiting until the current threads in kernel code can be preempted
 - This should be shorter than waiting for the current event to complete
 - Wait until lower priority threads release needed resources
 - This is a challenge for developers that must understand this requirement
 - Mastering synchronization becomes key to the overall correctness
 - Wait until higher priority tasks finish or block


Case Study – Unix Family of Operating Systems

15

- A traditional view of an operating system
 - Goes all the way back to the 70s
 - Promoting a certain programming model (philosophy even)
- Why study the Unix family?
 - Look into the motivation for the key concepts
 - Discuss the evolution of those concepts



- Early key concepts
 - A shared kernel and user-mode processes
 - Promoting C with the libC wrapping the ABI
 - Malloc and free for memory management
 - Streams (files and pipes)
 - Signals for traps and inter-process communication
 - Introduced device classes
 - All devices must fit into either a character or a block devices
 - With special files under /dev giving access to devices



- Device classes

- Character devices

- Stream of "characters"
 - Essentially serial lines
 - Examples: keyboards and terminals

Characters are bytes in fact (uint8_t)
Remember back then, ASCII codes where 7bits
No other languages than english were supported
Java popularized the use of Unicode (1995)

- Block devices

- Fixed-size blocks
 - Read/Write/Seek operations
 - Essentially tapes and disks

Network cards did not exist
Later forced into the block model
Now recognized as a new device type
Video cards did not exist
No concept in the kernel for 30 years
Changed with DRI

- Which programming model?

- Hardware is asynchronous
 - Thus event-oriented programming would be adequate
 - But stream-based programming was pushed forward
 - Then event-oriented programming was also adopted

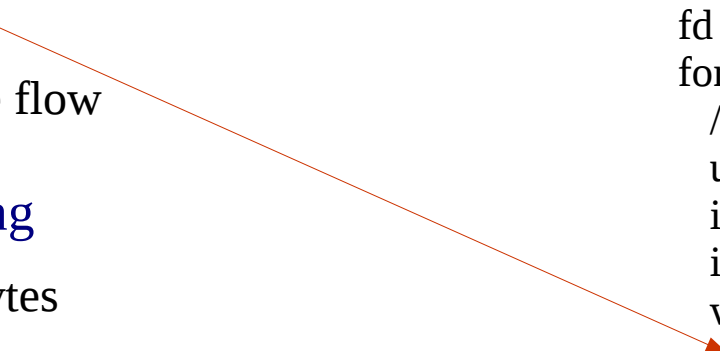
- Stream-based programming

- Open/read/write/close
- **Blocking operations**
- One algorithm as one flow

- Streams for everything

- Files as streams of bytes
- Pipes as communication channels
- Then sockets when networks appeared

```
void main(int argc, char* argv[]) {  
    int fd;  
    fd = open("toto");  
    for (;;) {  
        // read a record  
        uint8_t bytes[10];  
        int offset =0;  
        int remaining =10;  
        while (remaining) {  
            int n = read(fd, bytes+offset,remaining);  
            if (n==-1)  
                return;  
            offset += n;  
            remaining -= n;  
        }  
        // process record  
        ....  
    }  
}
```



```
struct read_event {
    struct event event;
    int fd;
    int offset, remaining;
    uint8_t *buffer;
#define STATE_INIT 0
#define STATE_READ 1
#define STATE_PROCESS 2
#define STATE_EOF 3
    uint8_t state;
};

void react(struct queue* rq, struct event* e) {
    struct read_event *evt = read_event_of(e);
    switch(evt → state) {
        case STATE_INIT:
            evt → fd = open("toto");
            evt → state = STATE_READ;
            event_push(rq,evt);
            break;
        case STATE_READ:
            read(evt)
            break;
        case STATE_PROCESS:
            process(rq,evt);
            break;
    }
}
```

Context Data


Finite State Machine

Non-blocking operation

```
void read(struct queue* rq,
          struct read_event* e) {
    while (remaining) {
        int n = read(fd, e → buffer+offset,remaining);
        if (n== -1) {
            evt → state = STATE_EOF;
            return;
        }
        if (n==0) {
            event_push(rq,e);
            return;
        }
        offset += n;
        remaining -= n;
    }
    evt → state = STATE_PROCESS;
    event_push(rq,e);
}
```

- Stream-based vs Event-based programming
 - A religious debate that is still going on
 - Stream-based programming
 - Yielded thread-based programming when threads appeared for multi-core
 - Event-oriented programming
 - GUI remained event-oriented, single threaded
 - The debate is hotter than ever on the server side
 - High performance servers started advocating multi-threading
 - Recently, servers are going towards event-oriented programming (e.g. node.js)
 - Embedded systems
 - A mix of both, some advocate threads, others defend events...
 - Android
 - A clear mix & a clear separation of concerns
 - Event-oriented for the GUI – All events run to completion
 - Threads for workers – Using blocking calls is allowed

- Devices as special files under /dev
 - As special files (mknod, major, minor)
 - Open/Close/Read/Write and optional Seek
 - Major-minor identifies a kernel module
 - Usually a device driver, but it could be any module



Require file system support
in the Kernel

```
$ ls -al /dev
C rw-r----- 1 root kmem    1,1  mem
C rw-rw-rw-  1 root root    1,3  null
C rw-rw-rw-  1 root root    1,5  zero

B rw-rw----  1 root disk    8,0  sda
B rw-rw----  1 root disk    8,1  sda1

C rw--w----  1 root tty     4,0  tty0
C rw-rw----  1 root tty     4,1  tty1

C rw-----  1 root root   254,0 rtc0

C rw-----  1 root root   10,239 uhid
C rw-----  1 root root   10,223 uinput
```

Also go have a look at...

/sys

/proc

Internal kernel data structures
reified through files

- Special files – the limitations of a file model


- The need for ioctl beyond read/write/seek
 - The man page says
 - “The `ioctl()` function manipulates the underlying device parameters of special files”
 - But in reality
 - It is a generic “function call” to the underlying device driver
 - Usually passing a buffer, whose content is described by the request parameter

```
#include <sys/ioctl.h>
int ioctl(int desc, unsigned long request, ...);
```

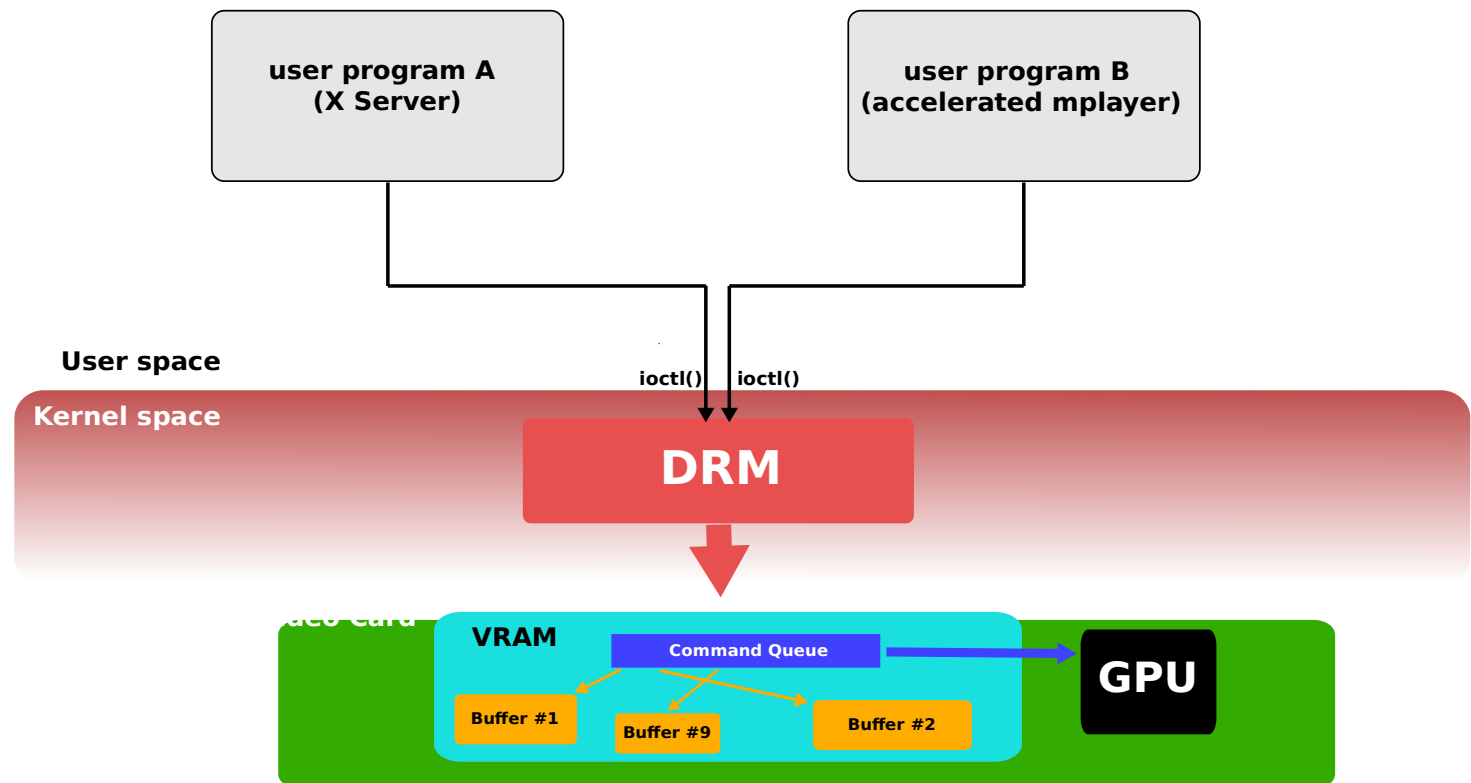
- The need to memory map special files
 - Best example is the frame buffer of the video card

```
#include <sys/mman.h>
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);
int munmap(void *addr, size_t length);
```

- An example – X11 Window Manager
 - Uses /dev to access devices
 - Because until very recently, the kernel had no concepts of a GUI
- Pointer device
 - Reads /dev/mouse to grab mouse events
 - Use ioctl to configure the mouse, if the mouse supports it
- Video device
 - Reads/writes to /dev/fb0 to display pixels onto the screen
 - Mmap the video buffer
 - Use ioctl to configure the video card
 - May use ioctl to access some accelerator functions
 - May use write to request accelerator functions and read for getting at the result
 - May also use the mmap memory with a circular buffer to send commands

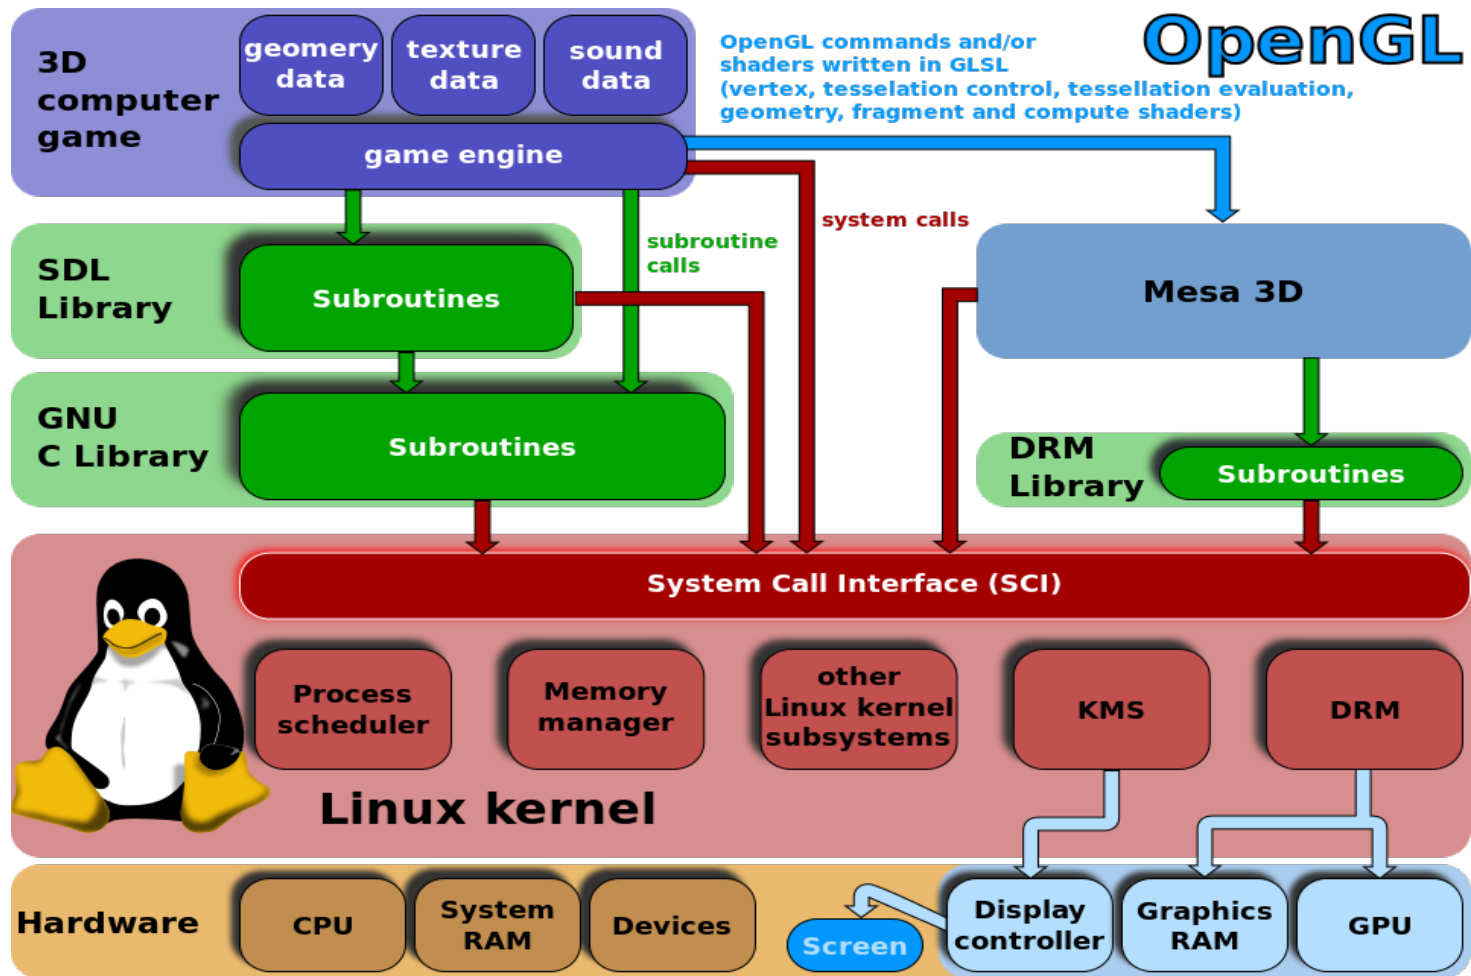


- Direct Rendering Infrastructure (DRI)
 - Direct Rendering Manager

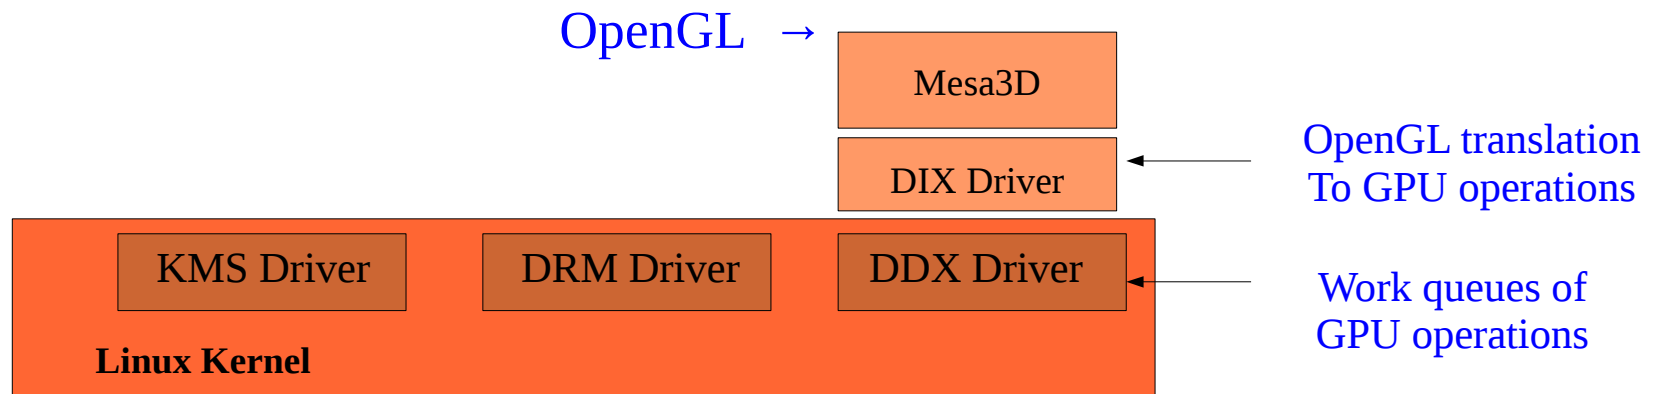


© 2014 Javier Cantero - this work is under the Creative Commons Attribution ShareAlike 4.0 license

- Direct Rendering Infrastructure (DRI)



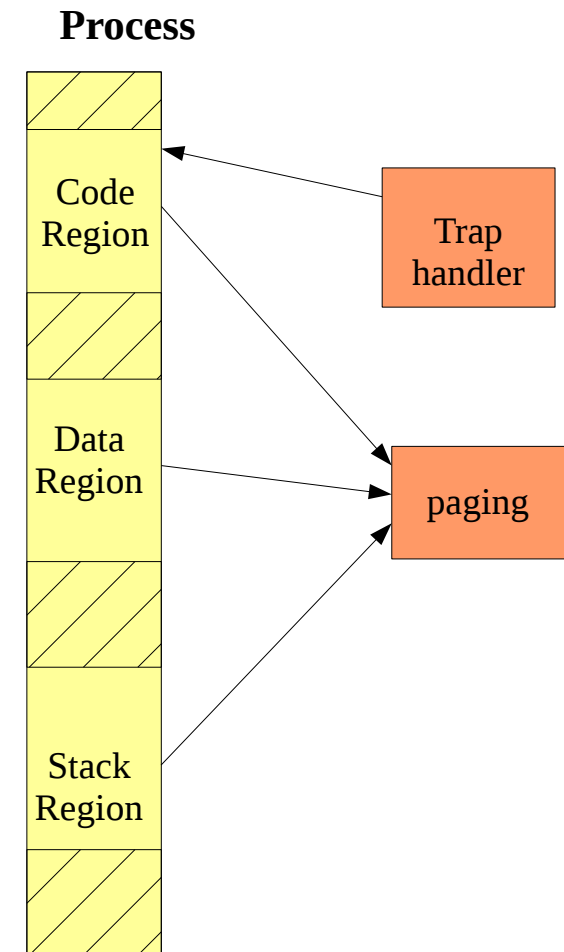
- Mesa3D
 - Open-source OpenGL
 - Simplified overview...



- Basic idea
 - A process provides a virtual machine
 - So that a real machine can be shared by multiple applications
- Process rationales
 - Memory isolation (safety, separation of concerns)
 - Security (access rights)
 - Paging (virtual memory larger than physical)
 - Management concept (kill, resource accounting)


- Processes as address spaces
 - A virtualized memory, isolated by hardware
 - Trap handler generates signals
- Mapped memory regions
 - Paging memory regions

➔ Paging requires
file system support
in the Kernel



- **Process management**
 - Create by forking (duplicate process)
 - Kill by pid
 - Exec a file → requires loaders
- **Loaders in the kernel**
 - Such as Elf loader or shell script loader
 - Elf loader
 - Reads in code/data section in code/data regions
 - Scrip loader
 - Loads the interpreter in (like /bin/bash)
 - Pass the script file as an argument

➡ Loaders require
file system support
in the Kernel



- Before threads
 - Every “task” was a process
 - Processes are heavy weight to create/destroy
 - Context switch between processes is expensive
 - Blocking calls meant a process was carrying a single task
- Original motivation
 - Light-weight multi-tasking within a process
- Example:
 - A server with multiple clients
 - One socket per client
 - One thread per client

**Preserves the programmin model
with blocking operations on streams**

- Multiple threads introduced the need for synchronization
 - Introduced mutexes and semaphores
 - Concurrent programming proved to be hard for most developers
 - Concurrency bugs are hard to find
 - Testing is much harder
- Multi-core evolution
 - Single core performance reached a peak early 2000
 - Multi-threading became also a performance quest
 - But beware, using threads is no guarantee of higher performance

So, do not confuse the two usages for threads

- **light-weight multi tasking, preserving the blocking programming model**
- **searching for high performance computing on multi-core machines**


- Multi-threading on multi-core
 - Synchronization has a cost in itself
 - Worst when using a syscall for each synchronization operation
 - User-level synchronization is faster, but strains the hardware (CAS-like typically)
 - Synchronization may induce convoy effects, drastically reducing parallelism
 - This happens on shared data structure
 - Context switch and cache efficiency
 - When context-switching between threads, the working set likely changes
 - Load balancing
 - Trade off between CPU usage and L1/L2 cache effectiveness

- Events are also supported
 - Almost all the ABI can be configured to be non-blocking
 - This allows to look at a thread as an event executor
 - It is possible to halt threads, waiting for events
 - With epoll for example, allows to wait on events on several sockets
- Graphical toolkits
 - Always promoted an event-oriented programming model
 - Some toolkits are single threaded, others permit the use of multiple threads


So, you have three programming models really, that can be combined

- **light-weight multi tasking, preserving the blocking programming model**
- **event-oriented programming using one or more threads**
- **using threads for high performance computing on multi-core machines**

- Boot process
 - From the hardware wakeup
 - Up to your application
 - How does it happen?
- Major steps
 - Hardware initialization
 - Boot loader
 - Kernel initialization
 - Distribution initialization




Memory



- Major steps

- Hardware initialization
- Boot loader
- Kernel initialization
- Distribution initialization

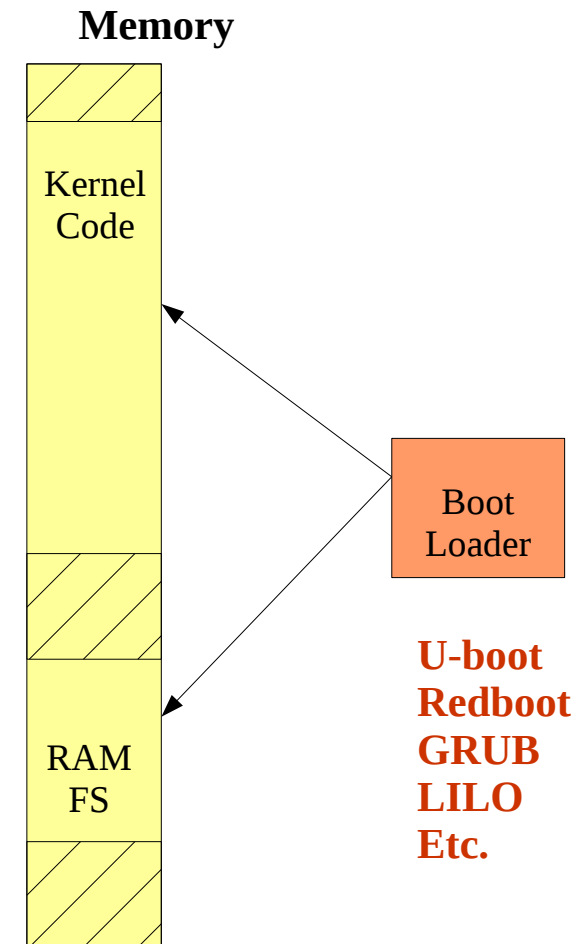


Nota Bene:

The boot loader needs access to many devices


- disks
- cd-rom or dvd
- USB

So it needs device drivers



- Major steps

- Hardware initialization
- Boot loader
- Kernel initialization
- Distribution initialization




Kernel self-initialize


Mounts the RAM file system as root "/"

Loads device drivers from the RAMFS

Create the first process from /init



- Major steps
 - Hardware initialization
 - Boot loader
 - Kernel initialization
 - Distribution initialization



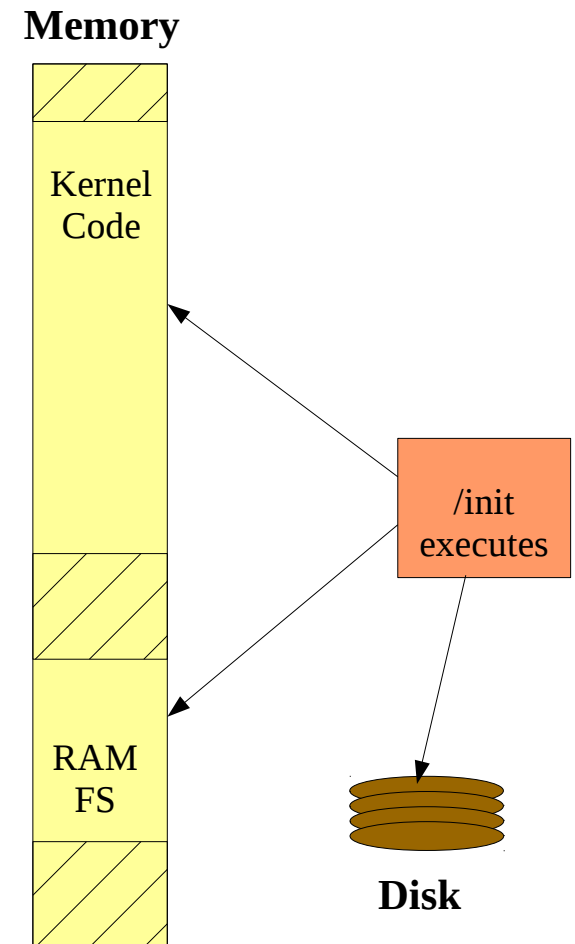
After that, it is a distribution initialization...

But most start with a shell /init

It mounts the real "/" from your disk

Then the root file system pivot happens

The rest of the initialization happens from your root partition



- Linux kernel vs Linux distribution
 - Two very different things
 - One kernel... (multiple versions of course)
 - Many distribution... with very different goals
 - Ubuntu, Debian, Raspbian, uc-linux, ...
 - Promote the C programming language and shell scripting
- Google Android
 - Linux kernel
 - A few core libraries (webkit, gstreamer, etc.)
 - Promote the Java programming language
 - But a Java-based distribution with its own software concepts
 - Components and services **distributed across processes**

- Unix-like operating systems
 - Use private files or communicate through pipes/sockets (safe programming)
 - Single threaded and blocking libC (easy programming)
 - File system crashes meant losing files and their content
- Opposite to DataBase Management Systems (DBMS)
 - Multi-threaded by essence, with shared data, protected through transactions
 - Strong query capabilities
 - Protected data, through failure recovery

You must reflect on these design differences!

Still true today?

Absolutely...

- OSGi Platforms

- Set-top boxes, ADSL boxes, factory-management gateways, etc.
- A world of services and modules
- Remote management – install/uninstall/start/stop
- But **everything runs in a single** Java Virtual Machine
- Runs on bare metal or on embedded operating system
- Also runs on Linux/Windows/MacOS

- Eclipse Platform

- OSGi platform
- GUI integration
- Resource integration

- Eclipse vs Linux

- Both virtual machines, both application platforms
- Both can run bare metal

✓	Memory Management	✓
✓	Threads and concurrency control	✓
✓	File system and sockets	✓
✓	Manage applications	✓
✓	Security	✓
⊘	Memory Isolation	✓
✓	Window Manager	✓
✓	Web Integration	✓

Not in Oracle Java but available in Android... Was not available in Windows for many years...
Some have proposed memory isolation in Java and other high-level languages