Design of Experiments

Arnaud Legrand and Jean-Marc Vincent

Scientific Methodology and Performance Evaluation
M2R MOSIG, Grenoble, September-December 2015

Outline

There are two key concepts:

<div align="center">replication and randomization</div>

You replicate to increase reliability. You randomize to reduce bias.

<div align="center">

**If you replicate thoroughly and randomize properly,
you will not go far wrong.**

</div>

There are two key concepts:

<div align="center">replication and randomization</div>

You replicate to increase reliability. You randomize to reduce bias.

**If you replicate thoroughly and randomize properly,
you will not go far wrong.**

It doesn't matter if you cannot do your own advanced statistical analysis. If you designed your experiments properly, you may be able to find somebody to help you with the statistics.

If your experiments are not properly designed, then no matter how good you are at statistics, your experimental efforts will have been in vain.

**Not even the most sophisticated statistical analysis can turn a
flawed experiment into a good one.**

Other important concepts:
- Pseudo-replication
- Experimental vs. observational data

Select the problem to study

Clearly define the kind of system to study, the kind of phenomenon to observe (state, evolution of state through time), the kind of study to conduct (descriptive, exploratory, prediction, hypothesis testing, . . .)

This is quite important as the set of experiments to perform will be completely different when you are:

- studying the stabilization of a peer-to-peer algorithm under a high churn
- trying to compare various scheduling algorithms or code versions
- modeling the response time of a server under a workload close to the server saturation
- . . .

This step will help you to determine which kind of experiment design you should use.

Define the set of relevant *response*

The system under study is generally modeled through a black-box model:
- some output variable/response(y)
- some inputs are fully unknown
- some input variables $(x_1, \ldots, x_p)$ are controllable
- whereas some others $(z_1, \ldots, z_q)$ are uncontrollable

Controlable factors
$x_1 \quad \cdots \quad x_p$

Inputs $\longrightarrow$ System $\longrightarrow$ Output y

$z_1 \quad \cdots \quad z_q$
Uncontrolable factors

In our case, the response could be:
- the makespan of a scheduling algorithm
- the amount of messages exchanged in a peer-to-peer system
- the convergence time of distributed algorithm
- the average length of a random walk
- the amount of energy or of memory used

Some of these metrics are the result of complex aggregation of measurements so they should be carefully recorded to check their correctness

Determine the set of relevant *factors* or *variables*

The system under study is generally modeled though a black-box model:
- some output variable/response(y)
- some inputs are fully unknown
- some input variables $(x_1, \ldots, x_p)$ are controllable
- whereas some others $(z_1, \ldots, z_q)$ are uncontrollable

Controllable factors

$x_1 \quad \cdots \quad x_p$

Inputs → | System | → Output y

$z_1 \quad \cdots \quad z_q$

Uncontrolable factors

Typical controllable variables could be:
- the heuristic used (e.g., FIFO, HEFT, . . .)
- one of their parameters (e.g., replication factor, a threshold, . . .)
- the size of the platform
- the degree of heterogeneity
- the version of the compiler

Uncontrollable variables could be:
- temperature, humidity, moon phase, road surface conditions
- someone using the machine and interfering with the experiment

You should carefully record all the factors you can think of

The typical case studies defined in the first step could include:

- Determining which variables are most influential on the response y (factorial designs, screening designs, analysis of variance)
  - Allows to distinguish between primary factors whose influence on the response should be modeled and secondary factors whose impact should be averaged
  - Allows to determine whether some factors interact in the response
- Devise an analytical model of the response y as a function of the primary factors x (regression, lhs designs)
- Fit a an analytical model (regression, response surface methodology, optimal designs)
  - Can then be used to determine where to set the primary factors x so that response y is always close to a desired value or is minimized/maximized
- Determining where to set the primary factors x so that variability in response y is small i.e., so that the effect of uncontrollable variables $z_1, \ldots, z_q$ is minimized (robust designs, Taguchi designs)
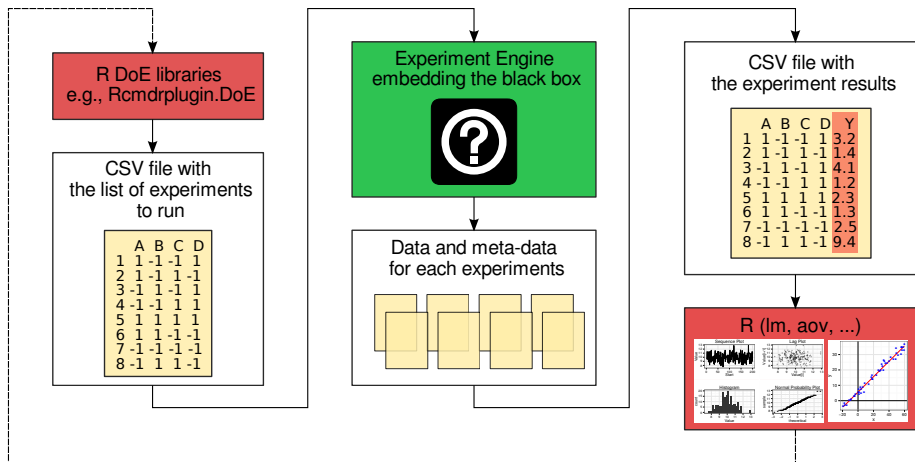
Outline

$$Y = a + bX + \varepsilon$$

- Y is the response variable
- X is a continuous explanatory variable
- a is the intercept
- b is the slope
- ε is some noise

When there are 2 explanatory variables:

$$Y = a + b^{(1)}X^{(1)} + b^{(2)}X^{(2)} + b^{(1,2)}X^{(1)}X^{(2)} + \varepsilon$$

ε is generally assumed to be independent of $X^{(k)}$, hence it needs to be checked once the regression is done

- Although your phenomenon is not linear, the linear model helps for initial investigations (as a first crude approximation)
- You should always wonder whether there is a way of looking at your problem where it is linear
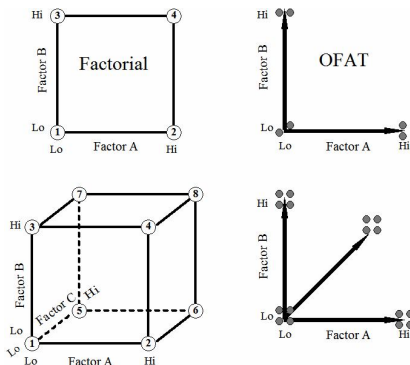
1. Decide a low and a high value for

The different values are by convention encoded with -1 and 1 but these are not *real* numbers

2. Test every (2^p) combination of high and low values, possibly replicating for each combination.
By varying everything, we can detect interactions right away

- Only a very small fraction of the space is covered (bias) ☹
- Interaction between factors cannot be estimated ☹
- Each replication allows to improve the estimation quality of only one factor, hence it requires more runs to have good estimates of all factors ☹

Unless dealing with a very simple problem, it is always better to change parameters all together than change parameters One Factor at a Time

```
1  library(FrF2)
2  d1 = FrF2(nruns=8 ,nfactors=3 , blocks=1 , replications = 2,
3  randomize= TRUE, seed= 26052 ,
4  factor.names=list(A=c(-1,1), B=c(-1,1), C=c(-1,1))); d1 ;
```

```
1    creating full factorial with 8 runs ...
2
3     run.no run.no.std.rp  A  B  C
4  1        1          2.1  1 -1 -1
5  2        2          6.1  1 -1  1
6  3        3          3.1 -1  1 -1
7  4        4          5.1 -1 -1  1
8  ...
9  15      15          1.2 -1 -1 -1
10 16      16          4.2  1  1 -1
11 class=design, type= full factorial
12 NOTE: columns run.no and run.no.std.rp are annotation, not part of t
```

How can we analyze something like this?

Outline

Confidence

If we had only 1 factor with 2 levels (2^1 design), the analysis would simply amount to compute confidence intervals or more precisely to test whether $\boxed{\mu_{A=Low} = \mu_{A=High}}$ or not (t-test)

(if few observations are available we would have to make the C.I wider and use the Student distribution)

But when having more factors and/or levels, we want to test whether some of the combinations have a significantly different expected value

Number of comparisons	2	3	4	5	6
Nominal Type I error	5%	5%	5%	5%	5%
Actual overall Type I error	5%	12.2%	20.3%	28.6%	36.6%

(See 16.1.5 of *Practical Regression and Anova using R* by Julian Faraway)

# Quick illustration of the difficulty of multiple testing

```
1  se = .1; mean = 7 ;
2  N = 10000;
3  df = data.frame(
4      x1=rnorm(N,mean=mean,sd=se),
5      x2=rnorm(N,mean=mean,sd=se),
6      x3=rnorm(N,mean=mean,sd=se))
7  df$eq1 = abs(df$x1-mean)<2*se
8  df$eq2 = abs(df$x1-df$x2)<2*se
9  df$eq3 = abs(df$x1-df$x2)<2*se &
10  abs(df$x1-df$x3)<2*se &
11  abs(df$x2-df$x3)<2*se
12
13  mean(df$eq1)
14  mean(df$eq2)
15  mean(df$eq3)
```

```
1  [1] 0.9538
2  [1] 0.8435
3  [1] 0.6596
```

ANOVA (ANalysis Of VAriance) enable to discriminate real effects from noise

- Enables to prove that some parameters have little influence and can be randomized over (possibly with a more elaborate model)
- Decomposes variance:

- Assumes identical standard deviation for the populations (homoscedastic)
- Multiple tests at once (assuming normality): $\mu_{A=Low,*} - \mu_{A=High,*} = 0$, $\mu_{B=Low,*} - \mu_{B=High,*} = 0$, ...

The ANOVA produces an F-statistic, the ratio of the variance calculated among the means to the variance within the samples.

- If the group means are drawn from populations with the same mean values, the variance between the group means should be lower than the variance of the samples
- A higher ratio therefore implies that the samples were drawn from populations with different mean values

The ANOVA produces an F-statistic, the ratio of the variance calculated among the means to the variance within the samples.

- If the group means are drawn from populations with the same mean values, the variance between the group means should be lower than the variance of the samples

- A higher ratio therefore implies that the samples were drawn from populations with different mean values

Let's work out a simple made-up example

```
1  Response = 10 + 2*as.numeric(d1$A) +
2      3*as.numeric(d1$B)*as.numeric(d1$C) + rnorm(nrow(d1))
3  d1 <- add.response(d1,Response, replace=TRUE)
```

I had to use `as.numeric` to interpret the -1 and 1 as numbers whereas they were created as factors

A simple ANOVA in R

```
d1_aov <- aov(Response ~ (A + B + C)^2, data=d1)
summary(d1_aov) # summary will call summary.aov
```

```
              Df Sum Sq Mean Sq  F value    Pr(>F)
A              1  22.98   22.98   38.318 0.000161 ***
B              1  68.02   68.02  113.417 2.11e-06 ***
C              1  77.60   77.60  129.402 1.21e-06 ***
A:B            1   0.44    0.44    0.728 0.415721
A:C            1   0.93    0.93    1.555 0.243804
B:C            1  14.62   14.62   24.374 0.000806 ***
Residuals      9   5.40    0.60
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

So, all factors are significant and there is a significant interaction between B
and C

Can't I just read my linear regression as usual?

```
1  summary.lm(d1_aov)
```

```
1  Call:
2  lm.default(formula = Response ~ (A + B + C)^2, data = d1)
3
4  Residuals:
5       Min       1Q    Median       3Q      Max
6  -1.01845 -0.48073 -0.01537  0.45886  0.98771
7
8  Coefficients:
9              Estimate Std. Error t value Pr(>|t|)
10 (Intercept)  19.5912     0.1936 101.194 4.56e-15 ***
11 A1            1.1984     0.1936   6.190 0.000161 ***
12 B1            2.0618     0.1936  10.650 2.11e-06 ***
13 C1            2.2023     0.1936  11.375 1.21e-06 ***
14 A1:B1         0.1652     0.1936   0.853 0.415721
15 A1:C1         0.2415     0.1936   1.247 0.243804
16 B1:C1         0.9558     0.1936   4.937 0.000806 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
19
20 Residual standard error: 0.7744 on 9 degrees of freedom
21 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
22 F-statistic:  51.3 on 6 and 9 DF,  p-value: 1.873e-06
```

```
1  summary.lm(d1_aov)
```

```
1  Call:
2  lm.default(formula = Response ~ (A + B + C)^2, data = d1)
3
4
5
6
7
8
9               Estimate Std. Error t value Pr(>|t|)
10 (Intercept)   19.5912     0.1936 101.194 4.56e-15 ***
11 A1             1.1984     0.1936   6.190 0.000161 ***
12 B1             2.0618     0.1936  10.650 2.11e-06 ***
13 C1             2.2023     0.1936  11.375 1.21e-06 ***
14 A1:B1          0.1652     0.1936   0.853 0.415721
15 A1:C1          0.2415     0.1936   1.247 0.243804
16 B1:C1          0.9558     0.1936   4.937 0.000806 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
19
20 Residual standard error: 0.7744 on 9 degrees of freedom
21 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
22 F-statistic:  51.3 on 6 and 9 DF,  p-value: 1.873e-06
```

Wait, why is the formula so different?

$$10 + 2A + 3BC$$

```
1  summary.lm(d1_aov)
```

```
1  Call:
2  lm.default(formula = Response ~ (A + B + C)^2, data = d1)
3
4
5
6
7
8
9                Estimate Std. Error t value Pr(>|t|)
10 (Intercept)   19.5912      0.1936 101.194 4.56e-15 ***
11 A1             1.1984      0.1936   6.190 0.000161 ***
12 B1             2.0618      0.1936  10.650 2.11e-06 ***
13 C1             2.2023      0.1936  11.375 1.21e-06 ***
14 A1:B1          0.1652      0.1936   0.853 0.415721
15 A1:C1          0.2415      0.1936   1.247 0.243804
16 B1:C1          0.9558      0.1936   4.937 0.000806 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
19
20                                             eedom
21                                          0.9527
22                                          -06
```

Wait, why is the formula so different?

$$10 + 2A + 3BC$$

Because it treated the factors "-1" and "1" as 0 and 1...

```
1 d1_lm <- lm(Response ~ (as.numeric(A) + as.numeric(B) +
2     as.numeric(C))^2, data=d1)
3 summary.aov(d1_lm)
```

```
1                          Df Sum Sq Mean Sq F value   Pr(>F)
2 as.numeric(A)             1  22.98   22.98  38.318 0.000161 ***
3 as.numeric(B)             1  68.02   68.02 113.417 2.11e-06 ***
4 as.numeric(C)             1  77.60   77.60 129.402 1.21e-06 ***
5 as.numeric(A):as.numeric(B) 1   0.44    0.44   0.728 0.415721
6 as.numeric(A):as.numeric(C) 1   0.93    0.93   1.555 0.243804
7 as.numeric(B):as.numeric(C) 1  14.62   14.62  24.374 0.000806 ***
8 Residuals                 9   5.40    0.60
9 ---
10 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Sweet, it's the same as the previous ANOVA

```
1 summary(d1_lm) # summary will call summary.lm
```

```
 1 Call:
 2 lm.default(formula = Response ~ (as.numeric(A) + as.numeric(B) +
 3     as.numeric(C))^2, data = d1)
 4
 5 Residuals:
 6      Min      1Q   Median      3Q     Max
 7 -1.01845 -0.48073 -0.01537  0.45886  0.98771
 8
 9 Coefficients:
10                             Estimate Std. Error t value Pr(>|t|)
11 (Intercept)                  15.4654     3.1870   4.853 0.000905 ***
12 as.numeric(A)                -0.0429     1.6878  -0.025 0.980277
13 as.numeric(B)                -2.6022     1.6878  -1.542 0.157516
14 as.numeric(C)                -2.7789     1.6878  -1.647 0.134064
15 as.numeric(A):as.numeric(B)   0.6606     0.7744   0.853 0.415721
16 as.numeric(A):as.numeric(C)   0.9658     0.7744   1.247 0.243804
17 as.numeric(B):as.numeric(C)   3.8232     0.7744   4.937 0.000806 ***
18 ---
19 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
20
21 Residual standard error: 0.7744 on 9 degrees of freedom
22 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
23 F-statistic:  51.3 on 6 and 9 DF,  p-value: 1.873e-06
```

```
1  summary(d1_lm) # summary will call summary.lm
```

```
1
2                                                      s.numeric(B) +
3
4
5
6
7
8
9
10                              Estimate Std. Error t value Pr(>|t|)
11 (Intercept)                   15.4654     3.1870   4.853 0.000905 ***
12 as.numeric(A)                 -0.0429     1.6878  -0.025 0.980277
13 as.numeric(B)                 -2.6022     1.6878  -1.542 0.157516
14 as.numeric(C)                 -2.7789     1.6878  -1.647 0.134064
15 as.numeric(A):as.numeric(B)    0.6606     0.7744   0.853 0.415721
16 as.numeric(A):as.numeric(C)    0.9658     0.7744   1.247 0.243804
17 as.numeric(B):as.numeric(C)    3.8232     0.7744   4.937 0.000806 ***
18 ---
19 Signif. codes:   0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
20
21 Residual standard error: 0.7744 on 9 degrees of freedom
22 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
23 F-statistic:  51.3 on 6 and 9 DF,  p-value: 1.873e-06
```
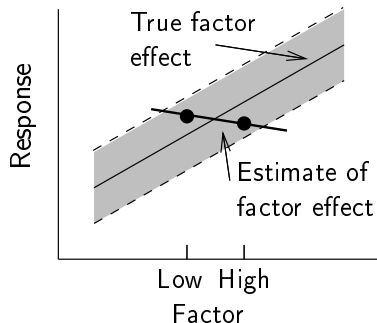
Variability is too large too obtain good estimates of the true coefficients

$$10 + 2A + 3BC$$

One should anyway use other kind of designs to estimate continuous model parameters

- The coding numbers are completely meaningless and influence the estimates of the slope
  - If your input parameters are numerical, go for *extreme* values, hoping the intermediate behavior is not too complicated and consider them as factors
- Real question: is a there significant increase when changing factors?
- Remember: you should use ANOVA for factorial designs, not LM
  - So don't use `summary.lm` in such cases; use `summary.aov`

```
1 MEPlot(d1, abbrev=4, select=c(1,2,3), response="Response")
```
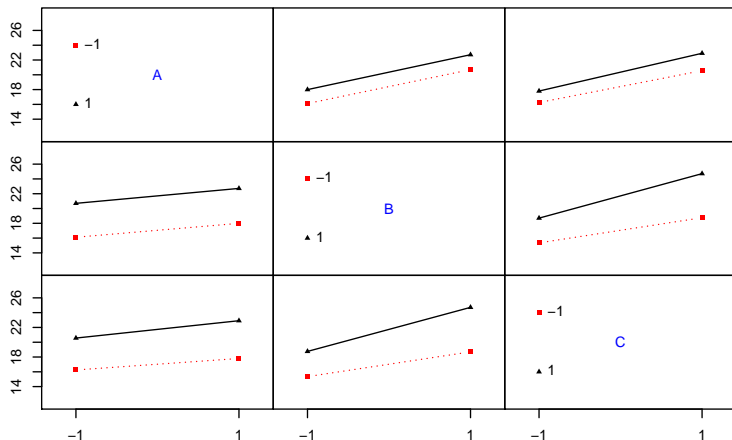

Main effects plot for Response

No CI on this one but we've seen that computing CIs is not straightforward
↝ rely on the `summary.aov`

What about interactions ?

```
1 IAPlot(d1, abbrev=4, show.alias=FALSE, select=c(1,2,3))
```


Interaction plot matrix for Response

Again, no CI $\rightsquigarrow$ rely on the `summary.aov`

Checking hypothesis

```
1 layout(matrix(c(1,2,3,4),2,2)) # optional layout
2 plot(aov(Response ~ (A + B + C)^2, data=d1))
```
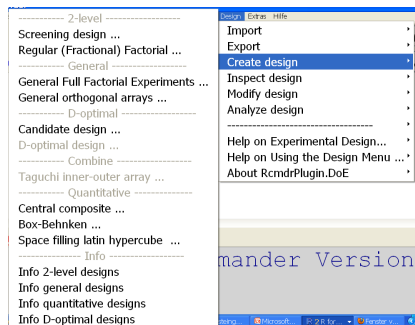
For the R commands, there is a trick: 😊

Use Rcmdr and RcmdrPlugin.DoE
(by Ulrike Grömping)

Simply `library(RcmdrPlugin.DoE)`...

You should only remember the principles and try to understand the underlying hypothesis

- ANOVA enables to discriminate real effects from noise in factorial experiments. *It relies on homoscedasticity and normality (or requires large number of samples)*
- 2-level factorial designs are a simple way to go and are more efficient than OFAT experiments
- Replicate thoroughly and randomize properly: you will not go far wrong

Outline

If $p = 8$, and the global variability is large, we may have to do $r = 5$ replications, hence $2^p.r = 256 \times 5 = 1280$ experiments!!!

- Then, you need something intermediate between OFAT and a full factorial 2^p design.
- It probably does not really make sense to study the joint effect of changing A, B, C, D, E, F, G, and H at the same time...

You should then go for a fractional 2^{p-k} design that will still make sure the combinations are well spread and the design is well balanced

Fractional designs

```
1  d2 = FrF2(nruns=8 ,nfactors=4 , blocks=1 , replications = 2,
2  randomize= TRUE, seed= 26052 ,
3  factor.names=list(A=c(-1,1), B=c(-1,1), C=c(-1,1), D=c(-1,1)));
4  d2 ;
```

```
1      run.no run.no.std.rp   A   B   C   D
2  1        1            2.1   1  -1  -1   1
3  2        2            6.1   1  -1   1  -1
4  3        3            3.1  -1   1  -1   1
5  4        4            5.1  -1  -1   1   1
6  5        5            8.1   1   1   1   1
7  ...
8  13      13            2.2   1  -1  -1   1
9  14      14            5.2  -1  -1   1   1
10 15      15            1.2  -1  -1  -1  -1
11 16      16            4.2   1   1  -1  -1
12 class=design, type= FrF2
13 NOTE: columns run.no and run.no.std.rp are annotation,
14       not part of the data frame
```

Not much gain here... Fractional designs have constraints but allow you to control how much you loose

Saving a lot of time/money: Plackett-Burman screening designs

```
1  d3 <-  pb(nruns= 20 ,n12.taguchi= FALSE ,nfactors= 20 -1, ncenter= 0 ,
2    replications= 1 ,repeat.only= FALSE ,randomize= TRUE ,seed= 26654 ,
3    factor.names=list( A=c(-1,1),B=c(-1,1),C=c(-1,1),D=c(-1,1),
4    E=c(-1,1),F=c(-1,1),G=c(-1, 1),H=c(-1,1),J=c(-1,1),K=c(-1,1),
5    L=c(-1,1),M=c(-1,1),N=c(-1,1),O=c(-1,1),P=c(-1,1) ) ) ; d3
```

	A	B	C	D	E	F	G	H	J	K	L	M	N	O	P	e1	e2	e3	e4
1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1	1
2	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1	1	1	1	1	-1	1
3	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
4	1	-1	-1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	1	-1	1
5	-1	1	1	1	-1	1	-1	1	1	-1	1	-1	-1	1	1	-1	1	-1	-1
6	-1	-1	1	1	1	-1	1	1	1	1	1	1	-1	-1	-1	1	-1	1	-1
7	1	1	-1	1	1	1	1	-1	1	1	1	-1	1	1	-1	-1	-1	-1	-1
8	-1	1	1	-1	1	1	1	1	-1	1	1	-1	1	-1	-1	-1	-1	1	-1
9	1	-1	1	-1	1	-1	1	1	-1	1	-1	1	1	1	1	1	1	-1	-1
10	1	1	-1	-1	1	-1	1	1	-1	1	1	-1	-1	-1	1	1	-1	-1	-1
11	-1	1	1	1	-1	1	1	-1	1	-1	1	1	1	-1	1	-1	-1	1	-1
12	1	1	-1	1	-1	1	-1	1	1	1	-1	1	1	-1	-1	1	-1	1	-1
13	-1	-1	1	1	1	1	-1	1	1	-1	1	1	-1	1	1	-1	1	1	-1
14	-1	1	-1	-1	-1	1	1	1	-1	1	-1	-1	1	1	1	1	-1	1	-1
15	1	-1	1	-1	-1	1	1	1	-1	-1	1	1	1	1	-1	-1	-1	1	1
16	1	1	-1	1	-1	1	-1	-1	1	-1	1	1	-1	-1	1	1			

Only allows to estimate primary factors, not interations

Preliminary step for further investigation

Outline

What about having more than two levels?

Before even considering the generation, how would this analyzed?

- ANOVA still works and interpretation is OK when there are one (1-way ANOVA) or two (2-way ANOVA) factors (with several levels)

- Otherwise, it is a nightmare to analyze and you should decrease either the number of factors or the number of levels

In term of design, you can still go for all combinations

General Full Factorial Experiments

```
1  d4 <- fac.design(nfactors= 2 ,replications= 3 ,repeat.only= FALSE ,
2   blocks= 1 , randomize= TRUE ,seed= 17366 ,
3   nlevels=c( 3,5 ), factor.names=list(
4   Size=c("S","M","L"),Color=c("R","G","B","M","Y"))) ; d4
```

```
1    creating full factorial with 15 runs ...
2
3      run.no run.no.std.rp Size Color
4  1        1           6.1    L     G
5  2        2          10.1    S     M
6  3        3          12.1    L     M
7  4        4           2.1    M     R
8  5        5          11.1    M     M
9  6        6          14.1    M     Y
10 7        7           4.1    S     G
11 8        8           3.1    L     R
12 9        9          13.1    S     Y
13 10      10          15.1    L     Y
14 11      11           5.1    M     G
15 12      12           1.1    S     R
16 13      13           8.1    M     B
17 14      14           7.1    S     B
18 15      15           9.1    L     B
19 16      16           1.2    S     R
20 17      17           8.2    M     B
```

Reducing the size of such designs

You can still sample from it but the outcome is likely to be not well balanced
- $\rightsquigarrow$ the estimation may not be that good and probably quite biased because of this 😣

```
d4[sample(size=5,replace=FALSE,1:nrow(d4)),]
```

```
    Size Color
29    L     G
30    L     M
41    L     G
3     L     M
25    S     B
```

That's why you should try to reduce as much as possible the number of factors and of levels if you can

Outline

Without any information about the response

Then we should not favor a region over an other
- What about all combinations of a regular division?

```
1  x <- seq(10, 100, length.out = 10)
2  y <- seq(0, 1, length.out = 10)
3  d5_regular <- expand.grid(A = x, B = y)
4  plot(d5_regular, main="Regular division")
```
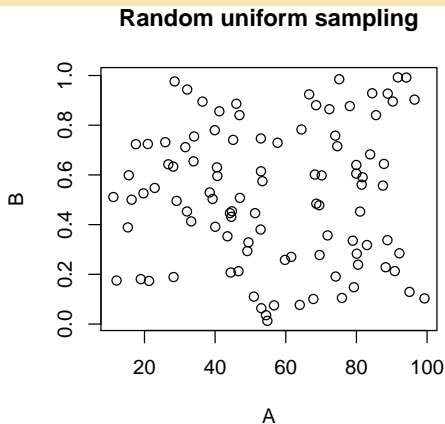

Regular division
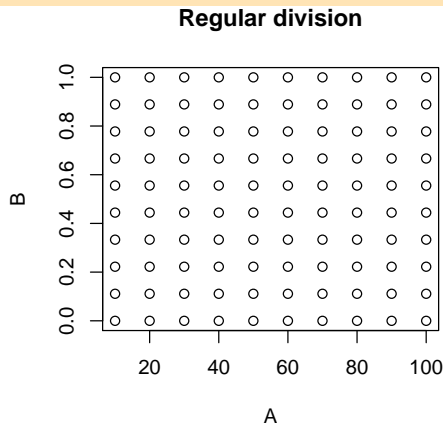
Can we have a less biased design?

We should not favor any particular value
- What about a uniform sampling then?

```r
1  set.seed(1);
2  x <- runif(100,min=10,max=100); y <- runif(100, min=0,max=1)
3  d5_unif <- data.frame(A = x, B = y)
4  plot(d5_unif, main="Random uniform sampling")
```


Random uniform sampling

Can we have a design covering better the whole space?

We do not want to miss any region
- Space filling designs: Latin Hyper Square designs and the maximin criteria

```
library(DoE.wrapper)
d5_maximin <- lhs.design( type= "maximin" , nruns= 100 ,nfactors= 2 ,
  digits= NULL , seed= 27041 , factor.names=list( A=c(10,100),B=c(0,1) ) )
plot(d5_maximin , select = c( "A","B" ), main="LHS design")
```


LHS design

This still reasonably works in higher dimensions

```r
library(DoE.wrapper);    set.seed(42);
d5_HD = lhs.design( type= "maximin" , nruns= 100 ,nfactors= 3 ,
     seed= 42 , factor.names=list( A=c(0,1),B=c(0,1),C=c(0,1) ) )
Response5 = 10 + 2*as.numeric(d5_HD$A) + 3*as.numeric(d5_HD$B)*as.numeric(d5_HD$C)
     rnorm(nrow(d5_HD),sd=1)
d5_HD <- add.response(d5_HD, Response5, replace=TRUE)
plot(d5_HD , select = c( "A","B","C" ), main="LHS design")
```


LHS design

What about the analysis?

```
1  summary(lm(Response5 ~ (A + B + C)^2, data = d5_HD))
```

```
1  Call:
2  lm.default(formula = Response5 ~ (A + B + C)^2, data = d5_HD)
3
4  Residuals:
5      Min      1Q   Median      3Q     Max
6  -2.90043 -0.64768  0.00095  0.75471  2.61620
7
8  Coefficients:
9              Estimate Std. Error t value Pr(>|t|)
10 (Intercept)  10.0932     0.5920  17.049   <2e-16 ***
11 A             1.5542     0.9686   1.605   0.1120
12 B             1.1188     0.8904   1.257   0.2121
13 C            -1.4085     0.9283  -1.517   0.1326
14 A:B          -2.3379     1.3228  -1.767   0.0804 .
15 A:C           3.0344     1.2428   2.442   0.0165 *
16 B:C           2.9668     1.2910   2.298   0.0238 *
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
19
20 Residual standard error: 1.087 on 93 degrees of freedom
21 Multiple R-squared:  0.451, Adjusted R-squared:  0.4156
22 F-statistic: 12.74 on 6 and 93 DF,  p-value: 1.909e-10
```

There is actually too much variability to conclude anything here (look at the R^2)

We know from the anova that B:C is significant but its Std. Error is still 1.29

We should add another round of 3 times more experiments to halve it

What happens if we fit a simpler model ?

```
1  summary(lm(Response5 ~ A + B:C, data = d5_HD))
```

```
1  Call:
2  lm.default(formula = Response5 ~ A + B:C, data = d5_HD)
3
4  Residuals:
5       Min        1Q    Median        3Q       Max
6  -3.00860  -0.71419  -0.00565   0.74843   2.98579
7
8  Coefficients:
9               Estimate Std. Error t value Pr(>|t|)
10 (Intercept)   10.0054     0.2471  40.489  < 2e-16 ***
11 A              1.8262     0.3920   4.659 1.01e-05 ***
12 B:C            3.0066     0.5247   5.730 1.13e-07 ***
13 ---
14 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
15
16 Residual standard error: 1.119 on 97 degrees of freedom
17 Multiple R-squared:  0.3938, Adjusted R-squared:  0.3814
18 F-statistic: 31.51 on 2 and 97 DF,  p-value: 2.852e-11
```

The Std. Errors decreased but remain quite high

As one could expect, the R^2 has decreased. . . 😣

```
1  Response5 = 10 + 2*as.numeric(d5_HD$A) +  3*as.numeric(d5_HD$B)*as.numeric(d5_HD$C)
2      rnorm(nrow(d5_HD),sd=.2) # Decreasing variability
3  d5_HD <- add.response(d5_HD, Response5, replace=TRUE)
4  summary(lm(Response5 ~ (A + B + C)^2, data = d5_HD))
```

```
1  Call:
2  lm.default(formula = Response5 ~ (A + B + C)^2, data = d5_HD)
3
4  Residuals:
5        Min       1Q   Median       3Q      Max
6  -0.50030 -0.10491 -0.00945  0.13446  0.47068
7
8  Coefficients:
9               Estimate Std. Error t value Pr(>|t|)
10 (Intercept) 10.06454    0.10992  91.558  < 2e-16 ***
11 A            1.58630    0.17986   8.820 6.41e-14 ***
12 B            0.13805    0.16533   0.835   0.4059
13 C            0.09524    0.17236   0.553   0.5819
14 A:B          0.46421    0.24562   1.890   0.0619 .
15 A:C          0.30745    0.23078   1.332   0.1860
16 B:C          2.33722    0.23972   9.750 6.92e-16 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
19
20 Residual standard error: 0.2019 on 93 degrees of freedom
```

One should actually instead fit the simple model suggested by the previous analysis:

$$y \sim A + B:C$$

You should use parsimony both in experiment design and modeling

```
1  summary(lm(Response5 ~ A + B:C, data = d5_HD))
```

```
1  Call:
2  lm.default(formula = Response5 ~ A + B:C, data = d5_HD)
3
4  Residuals:
5       Min        1Q    Median        3Q       Max
6  -0.56483  -0.11393   0.00626   0.12994   0.46614
7
8  Coefficients:
9              Estimate Std. Error t value Pr(>|t|)
10 (Intercept) 10.05536    0.04609  218.18   <2e-16 ***
11 A            1.94985    0.07311   26.67   <2e-16 ***
12 B:C          2.90476    0.09786   29.68   <2e-16 ***
13 ---
14 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
15
16 Residual standard error: 0.2087 on 97 degrees of freedom
17 Multiple R-squared:   0.95, Adjusted R-squared:  0.949
18 F-statistic: 921.8 on 2 and 97 DF,  p-value: < 2.2e-16
```

The principle of parsimony is attributed to the 14th century English philosopher William of Occam:

> *"Given a set of equally good explanations for a given phenomenon, the correct explanation is the simplest explanation"*

Parsimony (2/2)

The principle of parsimony is attributed to the 14th century English philosopher William of Occam:

> *"Given a set of equally good explanations for a given phenomenon, the correct explanation is the simplest explanation"*

- Models should have as few parameters as possible
- Linear models should be preferred to non-linear models
- Models should be pared down until they are *minimal adequate*

The principle of parsimony is attributed to the 14th century English philosopher William of Occam:

> *"Given a set of equally good explanations for a given phenomenon, the correct explanation is the simplest explanation"*

- Models should have as few parameters as possible
- Linear models should be preferred to non-linear models
- Models should be pared down until they are *minimal adequate*

This means, a variable should be retained in the model only if it causes a significant increase in deviance when removed from the current model

> *A model should be as simple as possible. But no simpler.*
>
> – *A. Einstein*

Outline

Working out a toy example

```r
x=lhs.design(type= "maximin", nruns=50, nfactors=1, seed=77,
      factor.names=list(x=c(0,60)))$x
y=3+x^2/60 + x*rnorm(length(x),sd=.3)
df = data.frame(x=x,y=y)
reg_quad <- lm(data=df,y~x+I(x^2))
summary(reg_quad)
```

```
Call:
lm.default(formula = y ~ x + I(x^2), data = df)

Residuals:
     Min       1Q   Median       3Q      Max
-21.7802  -4.5247   0.7544   5.1195  20.0284

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 2.124017   4.007473   0.530   0.5986
x           0.143694   0.310362   0.463   0.6455
I(x^2)      0.013169   0.005021   2.623   0.0117 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.483 on 47 degrees of freedom
Multiple R-squared:  0.7647, Adjusted R-squared:  0.7547
F-statistic: 76.36 on 2 and 47 DF,  p-value: 1.715e-15
```

We can clearly see where the heteroscedasticity comes from

```r
xv <- seq(0,60,.5)
yv <- predict(reg_quad,list(x=xv,x2=xv^2))
ggplot(data=df, aes(x=x,y=y)) + theme_bw() +
    geom_hline(yintercept=0) + geom_vline(xintercept=0) +
    geom_point(aes(x=x,y=y)) +
    geom_line(data=data.frame(x=xv,y=yv),aes(x=x,y=y),color="red")
```

Adding more points where there is more variability

```r
x=sqrt(lhs.design( type= "maximin" , nruns= 50 ,nfactors= 1 ,
        seed= 77 , factor.names=list( x=c(0,60^2) ) ) )$x)
y=3+x^2/60 + x*rnorm(length(x),sd=.3)

df = data.frame(x=x,y=y)
reg_quad <- lm(data=df,y~x+I(x^2))
summary(reg_quad)
```

```
Call:
lm.default(formula = y ~ x + I(x^2), data = df)

Residuals:
    Min      1Q  Median      3Q     Max
-27.256  -7.269   1.143   7.702  26.904

Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.292996  10.607678   0.028    0.978
x            0.257212   0.626398   0.411    0.683
I(x^2)       0.012031   0.008495   1.416    0.163

Residual standard error: 12.02 on 47 degrees of freedom
Multiple R-squared:  0.6569, Adjusted R-squared:  0.6423
F-statistic: 44.99 on 2 and 47 DF,  p-value: 1.209e-11
```


Histogram of x

```r
xv <- seq(0,60,.5)
yv <- predict(reg_quad,list(x=xv,x2=xv^2))
ggplot(data=df, aes(x=x,y=y)) + theme_bw() +
    geom_hline(yintercept=0) + geom_vline(xintercept=0) +
    geom_point(aes(x=x,y=y)) +
    geom_line(data=data.frame(x=xv,y=yv),aes(x=x,y=y),color="red")
```


The R^2 will never exceed 0.66 because our model fails fully explaining variance

- We should thus rather replicate for large values of x and average the results

- The expected value will be the same but the variance will be reduced
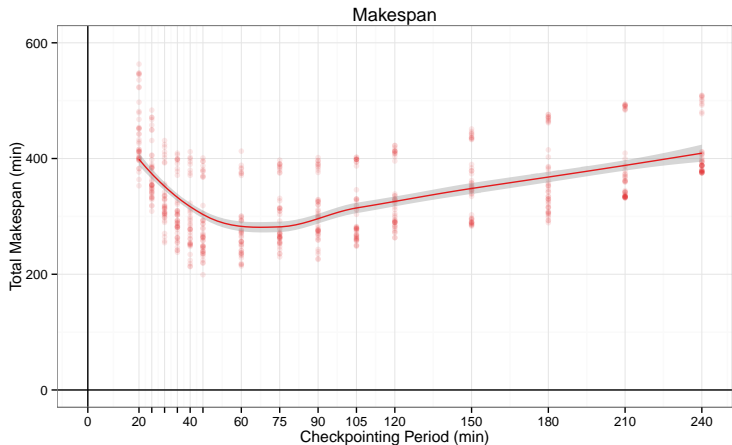
Outline

```
1  ggplot(data=df, aes(x=x,y=y)) + theme_bw() +
2      geom_hline(yintercept=0) + geom_vline(xintercept=0) +
3      geom_point(aes(x=x,y=y)) +
4      stat_smooth(method="loess",color="blue",fill="lightblue") +
5      geom_line(data=data.frame(x=xv,y=yv),aes(x=x,y=y),color="red") +
6      stat_function(fun=log) # the true function
```


LOcal RegrESSion: builds on linear regression to locally fit a line or a polynom

This is a very biased *estimator* so use with care

Discuss about the shape

Makespan

- "*the checkpointing period should be 68 minutes*": non-sense, uninteresting☹
- "*optimality region is flat and one should rather overestimate the checkpointing period*" 😀

Outline

When estimating model coefficients, it is intuitively better not to spread inputs but rather to use extreme values

- Note: this approach assumes that the model is correct

This intuitive notion can be formalized for linear models (see Hoos):

- Minimize generalized variance of least squares estimates of model parameters (determinant of covariance matrix)
  ⤳ D-optimal designs

- Minimize average variance (trace of covariance matrix)
  ⤳ A-optimal designs

- Minimize average of predicted response over experimental region
 ⤳ I-optimal designs

D-optimal Designs with R

```r
1  d7 <- lhs.design(type= "maximin", nruns= 200 , nfactors= 3,
2   digits=NULL, seed= 20521,
3   factor.names=list( A=c(0,1),B=c(0,1),C=c(0,1)))
4  d7.Dopt <- Dopt.design(25, data=d7, formula="~A +B:C", nRepeat= 5,
5        randomize= TRUE, seed=19583)
```

D-optimal Designs with R

```
1 d7 <- lhs.design(type= "maximin", nruns= 200 , nfactors= 3,
2  digits=NULL, seed= 20521,
3  factor.names=list( A=c(0,1),B=c(0,1),C=c(0,1)))
4 d7.Dopt <- Dopt.design(25, data=d7, formula="~A +B:C", nRepeat= 5,
5       randomize= TRUE, seed=19583)
```

- Designing experiments can be fun! 🙂
- Proceed carefully
  - The analysis is not simple but skilled statisticians can help you
  - The crucial part is actually the modeling, when you identify the factors, the response, and the kind of study
- This lecture only gives an overview but may already have changed your point of view on how to conduct experiments
- Remind the benefits of the sequential approach:
  - Parsimony
  - Use well-suited DoE and the corresponding analysis
  - Add measurements where there is variability

Recap on the lecture

1. Reproducibility is essential
   - literate programming with knitr or org-mode
   - laboratory notebook
2. Data manipulation and presentation
   - R, ggplot2, plyr, ...
3. Introduction to probabilities and statistics
   - A probabilistic model allows you to assess the confidence of your claims
4. Linear regression
   - The linear model is quite general
   - This knowledge about the system allows you to improve estimates
5. Design of Experiments
   - Sequential approach
   - Designs/analysis suited to every study